

Notebook

September 3, 2022

```
[1]: #PROYECTO FINAL TRATAMIENTO DE DATOS
```

```
##Francisco Salazar H.
```

```
[2]: import tensorflow as tf
import os
```

```
[3]: gpus = tf.config.experimental.list_physical_devices('GPU')
for gpu in gpus:
    tf.config.experimental.set_memory_growth(gpu, True)
```

```
[4]: tf.config.list_physical_devices('GPU')
```

```
[4]: []
```

```
[5]: import cv2
import imghdr
from matplotlib import pyplot as plt
```

```
[6]: data_dir = 'Data/Entrenamiento'
```


```
[7]: img = cv2.imread(os.path.join(data_dir, 'CLASS_02',
↳ '05-CAPTURE_20220421_053226_860.png' ))
```

```
[8]: img.shape
```

```
[8]: (216, 384, 3)
```

```
[9]: plt.imshow(img)
```

```
[9]: <matplotlib.image.AxesImage at 0x27b37486520>
```



```
[10]: #Load Data
```

```
[11]: import numpy as np
      from matplotlib import pyplot as plt
```


```
[12]: data = tf.keras.utils.image_dataset_from_directory(data_dir)
```

Found 275 files belonging to 2 classes.

```
[13]: data_iterator = data.as_numpy_iterator()
```

```
[14]: batch = data_iterator.next()
```

```
[51]: fig, ax = plt.subplots(ncols=4, figsize=(20,20))
      for idx, img in enumerate(batch[0][:4]):
          ax[idx].imshow(img.astype(int))
          ax[idx].title.set_text(batch[1][idx])
```



```

[52]: #Scale Data

[53]: data = data.map(lambda x,y: (x/255, y))

[54]: #Split Data

[55]: train_size = int(len(data)*.7)
      val_size = int(len(data)*.2)
      test_size = int(len(data)*.1)

[56]: train = data.take(train_size)
      val = data.skip(train_size).take(val_size)
      test = data.skip(train_size+val_size).take(test_size)

[57]: #Build a Deep Learning Model

[58]: train

[58]: <TakeDataset element_spec=(TensorSpec(shape=(None, 256, 256, 3),
dtype=tf.float32, name=None), TensorSpec(shape=(None,), dtype=tf.int32,
name=None))>

[59]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Flatten,
      ↪Dropout

[61]: model = Sequential()

[62]: model.add(Conv2D(64, (3,3), 1, activation='relu', input_shape=(256,256,3)))
      model.add(MaxPooling2D())
      model.add(Conv2D(32, (3,3), 1, activation='relu'))
      model.add(MaxPooling2D())
      model.add(Conv2D(16, (3,3), 1, activation='relu'))
      model.add(MaxPooling2D())
      model.add(Flatten())
      model.add(Dense(256, activation='relu'))
      model.add(Dense(1, activation='sigmoid'))

[63]: model.compile(loss='categorical_crossentropy',
                    metrics=['accuracy'])

[64]: model.summary()

```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 254, 254, 64)	1792

max_pooling2d_3 (MaxPooling 2D)	(None, 127, 127, 64)	0
conv2d_4 (Conv2D)	(None, 125, 125, 32)	18464
max_pooling2d_4 (MaxPooling 2D)	(None, 62, 62, 32)	0
conv2d_5 (Conv2D)	(None, 60, 60, 16)	4624
max_pooling2d_5 (MaxPooling 2D)	(None, 30, 30, 16)	0
flatten_1 (Flatten)	(None, 14400)	0
dense_2 (Dense)	(None, 256)	3686656
dense_3 (Dense)	(None, 1)	257

```

=====
Total params: 3,711,793
Trainable params: 3,711,793
Non-trainable params: 0
-----

```

```
[65]: #Train
```

```
[66]: logdir='logs'
```

```
[67]: tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=logdir)
```

```
[68]: train
```

```
[68]: <TakeDataset element_spec=(TensorSpec(shape=(None, 256, 256, 3),
dtype=tf.float32, name=None), TensorSpec(shape=(None,), dtype=tf.int32,
name=None))>
```

```
[69]: hist = model.fit(train, epochs=20, validation_data=val,
↳ callbacks=[tensorboard_callback])
```

Epoch 1/40

6/6 [=====] - 9s 1s/step - loss: 0.0000e+00 - accuracy: 0.3021 - val_loss: 0.0000e+00 - val_accuracy: 0.1875

Epoch 2/40

6/6 [=====] - 8s 1s/step - loss: 0.0000e+00 - accuracy: 0.2552 - val_loss: 0.0000e+00 - val_accuracy: 0.2188

Epoch 3/40

6/6 [=====] - 8s 1s/step - loss: 0.0000e+00 - accuracy: 0.2292 - val_loss: 0.0000e+00 - val_accuracy: 0.2188
Epoch 4/40
6/6 [=====] - 8s 1s/step - loss: 0.0000e+00 - accuracy: 0.2292 - val_loss: 0.0000e+00 - val_accuracy: 0.0938
Epoch 5/40
6/6 [=====] - 8s 1s/step - loss: 0.0000e+00 - accuracy: 0.2188 - val_loss: 0.0000e+00 - val_accuracy: 0.2188
Epoch 6/40
6/6 [=====] - 8s 1s/step - loss: 0.0000e+00 - accuracy: 0.2292 - val_loss: 0.0000e+00 - val_accuracy: 0.3125
Epoch 7/40
6/6 [=====] - 8s 1s/step - loss: 0.0000e+00 - accuracy: 0.2344 - val_loss: 0.0000e+00 - val_accuracy: 0.1250
Epoch 8/40
6/6 [=====] - 8s 1s/step - loss: 0.0000e+00 - accuracy: 0.2344 - val_loss: 0.0000e+00 - val_accuracy: 0.2500
Epoch 9/40
6/6 [=====] - 8s 1s/step - loss: 0.0000e+00 - accuracy: 0.2083 - val_loss: 0.0000e+00 - val_accuracy: 0.1250
Epoch 10/40
6/6 [=====] - 8s 1s/step - loss: 0.0000e+00 - accuracy: 0.2240 - val_loss: 0.0000e+00 - val_accuracy: 0.2188
Epoch 11/40
6/6 [=====] - 10s 2s/step - loss: 0.0000e+00 - accuracy: 0.1927 - val_loss: 0.0000e+00 - val_accuracy: 0.2812
Epoch 12/40
6/6 [=====] - 10s 2s/step - loss: 0.0000e+00 - accuracy: 0.2031 - val_loss: 0.0000e+00 - val_accuracy: 0.1562
Epoch 13/40
6/6 [=====] - 9s 1s/step - loss: 0.0000e+00 - accuracy: 0.1927 - val_loss: 0.0000e+00 - val_accuracy: 0.1562
Epoch 14/40
6/6 [=====] - 9s 1s/step - loss: 0.0000e+00 - accuracy: 0.2240 - val_loss: 0.0000e+00 - val_accuracy: 0.1562
Epoch 15/40
6/6 [=====] - 9s 1s/step - loss: 0.0000e+00 - accuracy: 0.2344 - val_loss: 0.0000e+00 - val_accuracy: 0.3125
Epoch 16/40
6/6 [=====] - 9s 1s/step - loss: 0.0000e+00 - accuracy: 0.2448 - val_loss: 0.0000e+00 - val_accuracy: 0.2500
Epoch 17/40
6/6 [=====] - 10s 2s/step - loss: 0.0000e+00 - accuracy: 0.2344 - val_loss: 0.0000e+00 - val_accuracy: 0.3438
Epoch 18/40
6/6 [=====] - 9s 2s/step - loss: 0.0000e+00 - accuracy: 0.2292 - val_loss: 0.0000e+00 - val_accuracy: 0.2812
Epoch 19/40

6/6 [=====] - 9s 1s/step - loss: 0.0000e+00 - accuracy: 0.2292 - val_loss: 0.0000e+00 - val_accuracy: 0.1562
Epoch 20/40
6/6 [=====] - 9s 1s/step - loss: 0.0000e+00 - accuracy: 0.2188 - val_loss: 0.0000e+00 - val_accuracy: 0.1875
Epoch 21/40
6/6 [=====] - 9s 2s/step - loss: 0.0000e+00 - accuracy: 0.2083 - val_loss: 0.0000e+00 - val_accuracy: 0.1562
Epoch 22/40
6/6 [=====] - 9s 2s/step - loss: 0.0000e+00 - accuracy: 0.2240 - val_loss: 0.0000e+00 - val_accuracy: 0.1875
Epoch 23/40
6/6 [=====] - 9s 1s/step - loss: 0.0000e+00 - accuracy: 0.2083 - val_loss: 0.0000e+00 - val_accuracy: 0.1875
Epoch 24/40
6/6 [=====] - 9s 2s/step - loss: 0.0000e+00 - accuracy: 0.2135 - val_loss: 0.0000e+00 - val_accuracy: 0.0938
Epoch 25/40
6/6 [=====] - 9s 2s/step - loss: 0.0000e+00 - accuracy: 0.2240 - val_loss: 0.0000e+00 - val_accuracy: 0.1562
Epoch 26/40
6/6 [=====] - 10s 2s/step - loss: 0.0000e+00 - accuracy: 0.2188 - val_loss: 0.0000e+00 - val_accuracy: 0.1562
Epoch 27/40
6/6 [=====] - 9s 2s/step - loss: 0.0000e+00 - accuracy: 0.1875 - val_loss: 0.0000e+00 - val_accuracy: 0.2500
Epoch 28/40
6/6 [=====] - 9s 2s/step - loss: 0.0000e+00 - accuracy: 0.2188 - val_loss: 0.0000e+00 - val_accuracy: 0.1250
Epoch 29/40
6/6 [=====] - 9s 1s/step - loss: 0.0000e+00 - accuracy: 0.2344 - val_loss: 0.0000e+00 - val_accuracy: 0.1875
Epoch 30/40
6/6 [=====] - 8s 1s/step - loss: 0.0000e+00 - accuracy: 0.2448 - val_loss: 0.0000e+00 - val_accuracy: 0.3125
Epoch 31/40
6/6 [=====] - 9s 1s/step - loss: 0.0000e+00 - accuracy: 0.2448 - val_loss: 0.0000e+00 - val_accuracy: 0.1250
Epoch 32/40
6/6 [=====] - 8s 1s/step - loss: 0.0000e+00 - accuracy: 0.2083 - val_loss: 0.0000e+00 - val_accuracy: 0.2500
Epoch 33/40
6/6 [=====] - 8s 1s/step - loss: 0.0000e+00 - accuracy: 0.2031 - val_loss: 0.0000e+00 - val_accuracy: 0.1250
Epoch 34/40
6/6 [=====] - 8s 1s/step - loss: 0.0000e+00 - accuracy: 0.2240 - val_loss: 0.0000e+00 - val_accuracy: 0.2188
Epoch 35/40

```

6/6 [=====] - 8s 1s/step - loss: 0.0000e+00 - accuracy:
0.2344 - val_loss: 0.0000e+00 - val_accuracy: 0.1250
Epoch 36/40
6/6 [=====] - 8s 1s/step - loss: 0.0000e+00 - accuracy:
0.2083 - val_loss: 0.0000e+00 - val_accuracy: 0.2812
Epoch 37/40
6/6 [=====] - 8s 1s/step - loss: 0.0000e+00 - accuracy:
0.2240 - val_loss: 0.0000e+00 - val_accuracy: 0.2188
Epoch 38/40
6/6 [=====] - 8s 1s/step - loss: 0.0000e+00 - accuracy:
0.2031 - val_loss: 0.0000e+00 - val_accuracy: 0.2812
Epoch 39/40
6/6 [=====] - 9s 2s/step - loss: 0.0000e+00 - accuracy:
0.2292 - val_loss: 0.0000e+00 - val_accuracy: 0.2500
Epoch 40/40
6/6 [=====] - 9s 2s/step - loss: 0.0000e+00 - accuracy:
0.2240 - val_loss: 0.0000e+00 - val_accuracy: 0.3125


```

```
[70]: #Plot Performance
```


```

[71]: fig = plt.figure()
plt.plot(hist.history['loss'], color='teal', label='loss')
plt.plot(hist.history['val_loss'], color='orange', label='val_loss')
fig.suptitle('Loss', fontsize=20)
plt.legend(loc="upper left")
plt.show()

```



```
[72]: fig = plt.figure()
plt.plot(hist.history['accuracy'], color='teal', label='accuracy')
plt.plot(hist.history['val_accuracy'], color='orange', label='val_accuracy')
fig.suptitle('Accuracy', fontsize=20)
plt.legend(loc="upper left")
plt.show()
```



```
[73]: #Evaluate
```

```
[74]: from tensorflow.keras.metrics import Precision, Recall, BinaryAccuracy
```

```
[75]: pre = Precision()
re = Recall()
acc = BinaryAccuracy()
```

```
[76]: for batch in test.as_numpy_iterator():
    X, y = batch
    yhat = model.predict(X)
    pre.update_state(y, yhat)
    re.update_state(y, yhat)
    acc.update_state(y, yhat)
```




```
[77]: print(pre.result(), re.result(), acc.result())
```


```
tf.Tensor(0.0, shape=(), dtype=float32) tf.Tensor(0.0, shape=(), dtype=float32)  
tf.Tensor(0.0, shape=(), dtype=float32)
```

```
[78]: #Test
```

```
[79]: img = cv2.imread('clase3.png')  
plt.imshow(img)  
plt.show()
```



```
[80]: resize = tf.image.resize(img, (256,256))  
plt.imshow(resize.numpy().astype(int))  
plt.show()
```



```
[81]: yhat = model.predict(np.expand_dims(resize/255, 0))
```

```
1/1 [=====] - 0s 68ms/step
```

```
[82]: yhat
```

```
[82]: array([[0.]], dtype=float32)
```

```
[84]: if yhat > 0.5:
      print(f'Clase predecida 2')
      else:
      print(f'Clase predecida 3')
```

```
Clase predecida 3
```